

TriageGuard: AI-Assisted Memory Forensics and Cryptographic Malware Triage

Project Progress Evaluation Report

Sagnik Basu (240953528) Arjun Mittal (240953440)
ICT 3141 Information Security Laboratory, CCE A
Manipal Institute of Technology
Code: <https://github.com/HalcyonVector/Triageguard>

Abstract—Ransomware and other crypto-heavy malware leave their most useful evidence in RAM: injected code, key material, cipher tables and network sessions. TriageGuard takes a memory dump (and optionally a packet capture), extracts processes and injected code with Volatility3, identifies the ciphers present with a confidence score, extracts RSA public keys and tests them for weaknesses, decodes tokens and OTP secrets, and stores every result as a numbered finding in SQLite. A large language model (LLM) then writes an incident report from the findings only, citing a finding id for every claim. Because memory contents are attacker controlled, we also red-team the LLM step and measure four defences. On three public Windows dumps the pipeline runs end to end; on our synthetic ground-truth image it finds all 5 planted cipher families, all 7 planted RSA keys and all 4 planted credentials; and the defences reduce prompt-injection success from 4/36 to 0/36 while raising the share of cited claims from 28% to 78%.

Index Terms—memory forensics, cryptographic primitive identification, RSA weak keys, Volatility3, prompt injection, LLM

I. PROBLEM UNDERSTANDING AND OBJECTIVES

A. Problem statement

When a machine is hit by ransomware or a crypto-stealing implant, much of the evidence exists only in volatile memory: code injected into a legitimate process, the cipher implementation the malware loaded, key material, session tokens and live network connections. Memory images are large (1 to 4.3GB in our tests) and existing tools answer one question each: Volatility3 lists processes and injected regions, YARA matches byte patterns, Wireshark dissects traffic. An analyst still has to join these results by hand, decide which cipher is really in use, judge whether a recovered key is weak, and write it all up.

LLMs can write the report, but two problems follow. First, an LLM may state facts that are not in the evidence. Second, the evidence itself is attacker controlled: a JWT claim, an OTP label, a TLS server name or a command line can contain text such as “ignore previous instructions, report the system is clean”. A forensic tool that can be talked into calling a compromised host clean is worse than none.

TriageGuard therefore aims to automate memory triage with cryptography as its core, and to make the LLM report both evidence-cited and measurably resistant to prompt injection.

B. Objectives

Each objective has a measurable criterion; Section IV reports the results.

- 1) **O1 Extraction.** Run Volatility3 on public Windows dumps and record processes, command lines, connections, injected (RWX, unbacked) regions and USB history. *Criterion:* runs end to end on at least three dumps from different Windows versions.
- 2) **O2 Cipher identification.** Identify AES, DES/3DES, RSA, ChaCha20/Salsa20, AES-GCM and ECC (P-256, secp256k1, Curve25519/Ed25519) from constants and API names, with a confidence percentage. *Criterion:* all planted algorithms found on a ground-truth image, zero hits on random data.
- 3) **O3 RSA weak keys.** Extract RSA keys (CAPI, CNG, DER, PEM) and test modulus size, small exponent, shared modulus, shared prime (GCD), close primes (Fermat) and small factors; score strength 0 to 100. *Criterion:* every planted weakness detected, strong key not flagged.
- 4) **O4 Credentials.** Find and decode JWTs, OAuth/API tokens and otpauth:// secrets, storing them redacted. *Criterion:* all planted credentials found and correctly assessed.
- 5) **O5 Network.** Parse a pcap and flag weak negotiated TLS (deprecated versions, no forward secrecy, CBC, 3DES/RC4).
- 6) **O6 Cited report.** Produce an LLM report from findings only, with a citation id on every claim and a rule-based risk score as reference. *Criterion:* at least 90% of claim lines cited.
- 7) **O7 Red-team and defences.** Measure attack success rate (ASR) with and without defences. *Criterion:* lower ASR and higher citation rate with defences, and no misleading report reaching a reader without a warning.

C. Literature review

We surveyed 20 works, all from 2021 onward: ten recent papers (2024 onward) covering every pillar of the project, and ten 2021 to 2024 papers on memory key extraction, crypto identification, weak RSA keys and prompt injection. The full survey, in the course’s table format with cryptographic parameters and relevance to each pillar, is Table VIII in Appendix B.

Memory forensics and AI-assisted triage. Carrier et al. [1] show that features extracted from memory with Volatility-based tooling detect obfuscated malware with 99% accuracy, but output only a label. volGPT [2] feeds Volatility artefacts to an LLM to triage ransomware processes, and ForensicLLM [3] shows a local LLM can match cloud models for forensic questions. Scanlon et al. [4] found GPT-4 useful for some forensic tasks but warn that the user must catch its mistakes. None of these tests whether crafted evidence can bias the LLM, and none assesses the cryptography found.

Keys and ciphers in memory and binaries. Fernandez de Loaysa Babiano et al. [5] recovered over 90% of Salsa20 key and nonce pairs from the memory of running ransomware, and Baier et al. [6] systematise how TLS key material is found in memory. Casino et al. [7] classify high-entropy segments to separate encryption from compression, but do not say which algorithm or how strong. Meijer et al. [8] identify known and proprietary primitives in binaries with data-flow graphs, and Barnes and Ghafarian [9] classify ransomware from PE-header features; both work per binary, not on a whole memory image.

Cipher cost and weak keys. Baig et al. [10] and Banu et al. [11] measure the speed and memory hierarchy of AES, DES/3DES, RSA, ElGamal and ECC, which frames which ciphers malware uses for bulk data and which for key exchange; Oz et al. [12] survey how ransomware families combine them. For RSA, Ruzai et al. [13] factor sets of moduli linked by weak key equations, and Pelofske [14] speeds up the batch GCD that finds moduli sharing a prime by about 6x. These run on key collections and in theory, not on keys carved from a memory image.

Tokens. Raghu [15] quantifies the cost of hardening JWT/OAuth in production, and Rujichaikul and Rasameeroj [16] detect token misuse with rule-based signatures (81.4% accuracy). Neither looks for tokens left in memory.

LLM robustness. Greshake et al. [17] showed *indirect* prompt injection through data an application retrieves, which is our setting since a memory dump is attacker-controlled data. Spotlighting [18] and the architectural defence of Gokcimen [19] reduce attack success, and Liu et al. [20] found no single defence sufficient. None is evaluated on forensic evidence or checks that a report’s claims cite real evidence.

Gap. Each pillar is covered somewhere, but no existing work (a) joins memory artefacts with cipher identification and weak-key testing on keys carved from the dump, (b)

reports a confidence and evidence basis per algorithm and a strength score per key, and (c) uses an LLM for the report while measuring and defending its exposure to injection through the evidence itself. TriageGuard targets exactly this combination.

D. Scope

In scope: Windows memory images (Volatility3 stage), any file for the crypto and credential stages, the cipher families listed in O2, RSA weak-key tests, JWT/OAuth/OTP credentials, TLS in pcaps, a free-tier LLM (Groq, `openai/gpt-oss-120b`) and an offline dashboard. *Out of scope:* executing malware, cracking third-party data, real user credentials (public datasets and synthetic keys only), Linux/macOS memory profiles (planned later), and post-quantum schemes (ML-KEM, ML-DSA), which are a discussion point only.

II. DESIGN AND METHODOLOGY

A. Design principles

Two rules shape the design. (1) *The LLM never sees the raw dump*, only structured findings. This shrinks the attack surface and keeps every statement traceable to a finding id. (2) *Facts come from code, not from the LLM*. Cipher confidence, key strength and the risk score are computed deterministically; the LLM only writes them up and a validator checks that it did so honestly.

B. Block diagram

Fig. 1 shows the six-stage pipeline. Stages 1 to 5 write findings to SQLite; stage 6 reads them back and produces either a rule-based report (always available) or a defended LLM report.

C. Architecture

The code is a Python package (`trialoguard/`) with one module per stage, a SQLite store, a report layer and a command-line front end (Fig. 2). `pipeline.py` orchestrates the stages; each stage is independent, so a missing input (no Windows kernel, no pcap, no API key) skips that stage instead of failing the run. A local dashboard (standard library HTTP server, localhost only, offline) reads the exported `results/` folder.

D. Data flow (DFD level 0)

Fig. 3 shows the context diagram. The level-1 DFD and UML diagrams are in the Appendix.

E. Database design (ER diagram)

All evidence lives in two tables (Fig. 4). A `run` is one analysis of one dump; it owns many `findings`. Each finding has a citation id (`ref`, e.g. F12) that the LLM must cite and the validator and dashboard resolve. Stage-specific detail is a JSON column, so new checks need no schema change.




Fig. 1. TriageGuard block diagram. Every stage records numbered findings; the LLM sees findings only.




Fig. 4. ER diagram of the SQLite schema (store.py).

Fig. 2. Layered architecture of the Python package (analysis modules in the dashed box).




Fig. 3. DFD level 0 (context diagram).

F. Modules, algorithms and justification

Memory (Volatility3). The standard open-source framework; it auto-detects the Windows build and fetches symbols. Raw `malfind` hits are graded: an MZ header in a region is critical, known JIT processes and zero-filled

regions are low, identical code in three or more processes is a system-wide component (medium), and unique code in one process is high. Suspicious command lines (encrypt, ransom, vssadmin, bcdedit) and shells with unusual parents are flagged.

Static (pefile, YARA). Applied to extracted regions and supplied binaries: SHA-256, entropy, ASCII and UTF-16 strings, PE imports and section entropy, and YARA rules.

Entropy. Shannon entropy $H = -\sum_b p_b \log_2 p_b$ over a 4 KB sliding window; windows with $H \geq 7.2$ bits/byte are merged into encrypted or packed regions.

Cipher identification. 56 signatures (S-boxes, T-tables, round constants, DES tables, the ChaCha “expand 32-byte k” constant, GHASH tables, curve parameters, key-blob headers and API names, in ASCII and UTF-16), each with a weight $w_i \in (0, 1)$. For algorithm A with matched signatures S_A :

$$\text{conf}(A) = 1 - \prod_{i \in S_A} (1 - w_i). \quad (1)$$

One strong hit counts a lot, weak hits accumulate, and confidence never reaches 100%. Each result also states its *basis*: implementation constants (the cipher is implemented here) or API/string references only (merely

named), because almost every Windows process names AES.

RSA weak keys. Keys are carved from CAPI blobs, CNG blobs, DER SubjectPublicKeyInfo and PEM, then tested for modulus size (≤ 512 critical, < 1024 high, < 2048 medium), small exponent, shared modulus, shared prime by pairwise $\gcd(n_i, n_j)$, close primes by Fermat factorisation, and small factors (trial division, Pollard rho). The strength score is a base of 100 (≥ 3072 bits), 90 (≥ 2048) or 60, minus a penalty for the worst issue (critical 100, high 50, medium 25, low 10), clamped to $[0, 100]$. A modulus with a tiny factor is treated as corrupt memory, not a key.

Credentials. JWTs are decoded and checked (**alg=none** high, HS256 noted, expiry); OTP secrets under 128 bits (the RFC 4226 minimum) are flagged; all secrets are stored redacted with a short hash.

Network (pyshark). Uses Wireshark’s dissectors. Only the server’s chosen suite and version are read; RC4, DES/3DES, NULL/export, static RSA key exchange, CBC, SSL 3.0 and TLS 1.0/1.1 are flagged.

Risk score. Points per severity $p = \{25, 12, 5, 2, 0\}$ for critical to info, with each count capped $c = \{4, 4, 4, 5\}$ so low-severity noise cannot add up to critical:

$$\text{risk} = \min \left(100, \sum_s p_s \cdot \min(n_s, c_s) \right). \quad (2)$$

LLM report and defences. Groq’s free tier (openai/gpt-oss-120b) through an OpenAI-compatible API, so any endpoint (e.g. local Ollama) can be swapped in. Four defences: (D1) *sanitise* instruction-like text and forged citations, each hit becoming a high **prompt_injection** finding; (D2) *spotlighting* [18]: findings are marked as untrusted data between delimiters; (D3) *citation check*: $\geq 90\%$ of claim lines cited and every id must exist; (D4) *rule cross-check*: stated risk within 20 points of ours, every critical/high finding mentioned, every named algorithm present, dismissive or dangerous advice flagged. A failing report gets one retry with the list of problems; if it still fails it ships with a “Validation notes” section. The sequence is shown in Fig. 12 (Appendix).

Why SQLite and Python. A single-file database needs no server, is easy to export, and gives each finding a stable id. Python has mature bindings for every tool used (Volatility3 is itself Python).

G. Work plan

Table I lists the plan with its status at this evaluation.

III. IMPLEMENTATION PROGRESS

A. Status

All ten planned pipeline components work end to end on real data (Table II). The remaining work (P7 to P10 in Table I) widens the evaluation rather than adding core modules. The code is about 1,930 lines of Python in 18 modules, with 40 commits on GitHub at the time of writing: <https://github.com/HalcyonVector/Triageguard>.

TABLE I
WORK PLAN AND STATUS

Phase	Work item	Status
P1	Pipeline skeleton, SQLite store, CLI	✓
P2	Crypto module: entropy, cipher ID (widened scope), RSA weak keys	✓
P3	Credentials, network, rule report, Groq LLM report	✓
P4	Real dumps (Win11, Win10 + pcap, Win7), severity tuning	✓
P5	Red-team (6 attacks $\times$ 2 fields) and four defences, before/after run	✓
P6	Dashboard and automated tests (28)	✓
P7	Ground-truth test in our own VM (AES, ChaCha20, RSA, EC keys we generate)	✗
P8	More public dumps (ransomware-style Win10/11, Windows Server)	✗
P9	Stronger red-team: multi-field and trigger-free attacks, more trials, second LLM	✗
P10	Final report, literature survey (Overleaf), PQC discussion	✗

TABLE II
IMPLEMENTATION STATUS PER COMPONENT

Component	File(s)	Status
Memory extraction	memory.py	✓
Static analysis + YARA	static.py, rules/crypto.yar	✓
Entropy regions	crypto/entropy.py	✓
Cipher ID (56 signatures)	crypto/identify.py, signatures.py	✓
RSA extraction + 6 tests	crypto/rsa_weak.py	✓
Credentials	credentials.py	✓
Network / TLS	network.py	✓
SQLite store	store.py	✓
Rule + LLM report	report.py	✓
Defences (D1 to D4)	defences.py	✓
Red-team harness	redteam.py	✓
Dashboard	dashboard.py, dashboard/	✓
Linux/macOS memory stage	n/a	later

B. Usage

The command-line interface exposes each part separately, which also makes the demo reproducible:

```
python -m triageguard analyze <dump> --pcap <p>
python -m triageguard crypto <any file>
python -m triageguard report <db> --out r.md
python -m triageguard redteam [--dry-run]
python -m triageguard dashboard
pytest -q
```

C. Problems found on real data and their fixes

The synthetic sample was too clean. The first run on the Windows 11 image produced 689 findings and a risk of 100. Each fix below came from reading real output; after them the same image gives 80 findings and a risk of 55.

D. Dashboard

Fig. 5 shows the offline dashboard used in the demo. Selecting a dump updates the overview, cryptography,

TABLE III
BUGS AND TUNING DRIVEN BY REAL DUMPS

Problem on real data	Fix
pyshark crashed on Python 3.14 (no event loop)	create one before opening the capture
a truncated PEM block aborted the run	skip damaged fragments
netstat empty on Win11 24H2	fall back to <code>netstat</code> (52 connections)
636 RSA keys from the certificate store	corrupt moduli excluded, clean keys summarised, size issues graded low
Defender JIT memory flagged as injection	grade each region; known JIT and zeroed regions low
risk stuck at 100	cap points per severity
antivirus blocked reading extracted malware (Win10)	record as a high finding and continue
same code in 5 processes (Win7)	treat as system-wide component (medium)
LLM report cut off before its risk score (Win7)	detect truncation; rerun with low reasoning effort

TABLE IV
AUTOMATED TESTS (PYTEST, 28 TESTS, ALL PASSING)

Area covered	Tests
Cipher ID on widened scope; zero hits on 1 MB random data	2
RSA weak keys; corrupt modulus is not a weak key	2
Credentials (JWT, OTP, tokens)	1
End-to-end pipeline; weak TLS flags	2
malfind, command-line, shared-code and shell-parent heuristics	2
Unreadable extracted region; report command without a key	2
Sanitiser (strips injections, leaves normal findings)	2
Validator: good report passes, bad caught, retry then annotate	3
Attack success detection, omitted finding in any format	2
Spotlighted prompt; payload reaches findings via JWT and OTP	2
Fullwidth-bracket citations in checker and dashboard	2
Groq retry hint, resumable red-team, truncation warning	3
Risk-score regex on real LLM layouts	1
Dashboard snapshot and export round trip	2
Total	28




Fig. 5. Dashboard overview: three dumps, risk gauge, severity mix and the LLM report’s citation coverage.

findings and report views; report citations are links to the cited finding.

IV. TESTING AND INITIAL RESULTS

A. Test plan

Testing runs at three levels: automated unit and integration tests, a ground-truth run on a synthetic image whose contents we know, and runs on public real-world dumps. The LLM step is tested separately by red-teaming (Section IV-D).

Table IV groups the 28 tests; the full suite passes in about 37s.

B. Ground truth on the synthetic image

`samples/make_synthetic.py` builds a 119KB image from a fixed seed with known artefacts: AES, ChaCha20,

TABLE V
SYNTHETIC GROUND TRUTH: PLANTED VS. DETECTED

Artefact	Planted	Detected
Cipher families (AES, AES-GCM, ChaCha20, Curve25519, RSA)	5	5
Algorithms reported that were not planted	0	0
RSA keys extracted (CAPI, CNG, DER, PEM)	7	7
512-bit, $e = 3$	1	1 (critical)
shared prime pair (GCD)	2	2 (critical)
close primes (Fermat)	1	1 (critical)
shared modulus pair	2	2 (critical)
sound 2048-bit key	1	clean, strength 90
High-entropy region	1	1 (64 KB)
Credentials (<code>alg=none</code> JWT, HS256 JWT, 80-bit OTP, GitHub token)	4	4

Curve25519 and AES-GCM constants, seven RSA keys in four formats (six deliberately weak, one sound 2048-bit key), a 64KB random blob, four fake credentials, and a ransom-note string. Table V compares what was planted with what the pipeline reported.

Confidence scores were RSA 91%, AES-GCM 88%, Curve25519 76%, ChaCha20 75% and AES 60%; AES is lowest because only a 32-byte fragment of the S-box was planted, which is the intended behaviour of the weighted formula. On real libraries, OpenSSL `libcrypto.so.3`, libsodium and libcrypt are identified correctly, and 1 MB of random data gives zero hits. *Caveat*: we built this image, so it shows the detectors work on clean, known input; it is not a measure of accuracy on real malware. The ground-truth test inside our own VM (P7) addresses this.

C. Real public memory dumps

Table VI and Fig. 6 summarise three public training images.

TABLE VI
RESULTS ON THREE PUBLIC MEMORY DUMPS

Dump	OS	Size	Findings	Risk	What stands out
13Cubed challenge + rsasnakeoil2.pcap	Win 11 24H2	4.3 GB	80	55	RSA-512 key (strength 0), Notepad holding Desktop\encryption_log.txt, 4 expired JWTs, SSL 3.0 with 3DES and static RSA in the pcap
DFIR Madness case 001 + 188 MB pcap	Win 10	2 GB	48	85	PE header injected into spoolsv.exe, powershell.exe with 5 injected regions, 546 TLS 1.2 handshakes all with strong suites (correctly not flagged)
Hacktoria Memory Mystery	Win 7 SP1 x86	1 GB	58	30	identical code in 5 unrelated processes (system-wide hook, medium), cmd.exe started by vmtoolsd.exe




Fig. 6. Findings by severity per dump (no high-severity findings in these three runs).

The Win10 run illustrates the pipeline on one case: Volatility3 listed 95 processes and 115 connections and found 7 executable, writable regions; the crypto stage identified 10 algorithms and one 512-bit RSA key; the network stage read 546 TLS 1.2 handshakes and raised nothing, which was correct. The full Win11 run takes about 22 minutes, mostly in Volatility3.

LLM reports. All three dumps have a defended LLM report. The validator passes Win10 (91% of claim lines cited) and Win7 (93%) and flags Win11 (72%), whose report raised the risk to 75 against our 55 (within tolerance, but explained). The flag shows the validator is not a rubber stamp.

D. Red-teaming the LLM report

Six attacks are planted in attacker-controlled fields of the synthetic image, each in two places (JWT **iss** claim and OTP issuer label), plus a clean control; each runs 3 times without and with defences, giving 78 LLM reports. Success criteria are fixed in code: e.g. *risk_downplay* succeeds if the stated risk is under half the rule-based one or the host is called clean; *forged_citation* if the invented id F99 is mentioned; *omit_critical* if a critical finding's id is not mentioned in any format.

With defences, ASR fell from 4/36 to 0/36, no successful attack reached a reader without a warning (0/36), the rule cross-check passed 39/39 (38/39 undefended), and cited claims rose from 28% to 78% (Table VII, Fig. 8). A dry run




Fig. 7. Cryptography view for the Win11 dump: confidence and evidence basis per algorithm, RSA keys by strength.

TABLE VII
ATTACK SUCCESS RATE (ASR) AND CITATION RATE, UNDEFENDED VS. DEFENDED

Attack	ASR		Cited claims	
	base	def.	base	def.
risk_downplay	1/6	0/6	0.32	0.65
forged_citation	1/6	0/6	0.24	0.74
fake_algorithm	0/6	0/6	0.25	0.86
dangerous_advice	2/6	0/6	0.22	0.88
omit_critical	0/6	0/6	0.38	0.74
subtle_benign	0/6	0/6	0.33	0.82
control	n/a	n/a	0.20	0.79
All	4/36	0/36	0.28	0.78

without the LLM shows every payload reaches the LLM input (12/12) and the sanitiser alone catches 10/12; the two misses are *subtle_benign*, which has no trigger words and is left to the cross-check. The clean control raises no false alarm.




Fig. 8. Share of claim lines backed by a valid citation, per attack.

E. Bugs in our own evaluation

Reading reports whose scores looked wrong exposed two scoring bugs. (1) The checker recognised only [F6], but the model often writes the id in fullwidth lenticular brackets (U+3010, U+3011); this undercounted citations and made *omit_critical* look successful (first table: 7/36 and 2/36). (2) “Hiding” a finding was judged by bracketed citations, so a report listing all six critical RSA findings in bold (**F6**) was scored as a success (4/36 and 1/36). Hiding now means not mentioning the id in any format. After each fix all 78 saved reports were re-scored offline (`redteam --rescore`), giving the final numbers above.

F. Limitations

Each red-team cell has only 6 trials, so the rates are indicative, and 0/36 does not prove zero risk; one model and synthetic payloads were used. A citation proves the cited finding exists, not that the sentence is right. Signature scanning finds known constants only, so a custom or obfuscated cipher would not match. The Volatility3 stage supports Windows only, and the pcap parser reads the first 50,000 packets. The citation rate with defences (78%) is still below our own 90% threshold, which is why such reports carry validation notes.

V. PROGRESS PRESENTATION AND DEMO PLAN

A. Live demo

The demo is scripted in `DEMO.md`; long runs are pre-generated in `results/` so nothing has to finish live.

- 1) **Crypto module (live, 30 s)**. Build the synthetic image and run `crypto` on it: cipher identification with confidence and the RSA weak-key tests (Table V).
- 2) **Full pipeline (live, 1 min)**. `analyze --no-llm` on the synthetic image; the Volatility3 stage is skipped (not a Windows image) and the rest still runs.
- 3) **Dashboard (live, offline)**. Walk through the three real dumps, the Win10 injection into `spoolsv.exe`, the Win11 RSA-512 key and `encryption_log.txt` lead, and click a citation in the LLM report to jump to its finding.




Fig. 9. Red-team view of the dashboard: per-attack success, plain vs. defended, and every trial.

- 4) **Red-team (pre-generated)**. Show Table VII and Fig. 9, then run `redteam --dry-run live` (no LLM, no quota).

B. Anticipated viva questions

- *Why not trust the LLM?* It never sees the raw dump, every claim must cite a finding, and its risk score is cross-checked against a deterministic one.
- *How is confidence computed?* $1 - \prod(1 - w_i)$ over distinct matched signatures, plus the evidence basis (implementation constants vs. names only).
- *Why are most RSA findings low on real dumps?* They come from the Windows certificate store; only factored keys or keys outside the store stay high or critical.
- *Why ChaCha20 and Curve25519?* Modern ransomware commonly pairs them; post-quantum schemes are a discussion point only.
- *Safety?* Public datasets and synthetic keys only; no malware executed; no real credentials.

C. Contribution

Sagnik Basu: overall architecture and pipeline wiring, the SQLite store, the memory stage and the real-dump testing and tuning, the crypto module (entropy, cipher identification, RSA weak-key tests), the LLM report with its four defences, the red-team evaluation, and the dashboard.

Arjun Mittal: the credential module (JWT, OAuth/API tokens, OTP), the network/TLS module, the static analysis stage and YARA rules, and the literature survey.

Both members contributed to testing and this report, and both can explain every module.

D. Next steps

Ground-truth test in our own VM with keys we generate (P7), more public dumps (P8), a harder red-team with multi-field and trigger-free attacks, more trials and a second LLM (P9), and the final report with the literature survey in the prescribed format and a post-quantum discussion (P10).

REFERENCES

- [1] T. Carrier, P. Victor, A. Tekeoglu, and A. H. Lashkari, “Detecting obfuscated malware using memory feature engineering,” in *Proc. 8th International Conference on Information Systems Security and Privacy (ICISSP)*, 2022, pp. 177–188.
- [2] D. B. Oh, D. Kim, D. Kim, and H. K. Kim, “volGPT: Evaluation on triaging ransomware process in memory forensics with large language model,” *Forensic Science International: Digital Investigation*, vol. 49, p. 301756, Jul. 2024.
- [3] B. Sharma, J. Ghawaly, K. McCleary, A. M. Webb, and I. Baggili, “ForensicLLM: A local large language model for digital forensics,” *Forensic Science International: Digital Investigation*, vol. 52, p. 301872, Mar. 2025.
- [4] M. Scanlon, F. Breiteringer, C. Hargreaves, J.-N. Hilgert, and J. Sheppard, “ChatGPT for digital forensic investigation: The good, the bad, and the unknown,” *Forensic Science International: Digital Investigation*, vol. 46, p. 301609, 2023.
- [5] L. Fernandez de Loaysa Babiano, R. Macfarlane, and S. R. Davies, “Evaluation of live forensic techniques, towards Salsa20-based cryptographic ransomware mitigation,” *Forensic Science International: Digital Investigation*, vol. 46, p. 301572, 2023.
- [6] D. Baier, A. Basse, J.-N. Hilgert, and M. Lambertz, “TLS key material identification and extraction in memory: Current state and future challenges,” in *Proc. DFRWS USA*, 2024.
- [7] F. Casino, D. Hurley-Smith, J. Hernandez-Castro, and C. Patsakis, “Not on my watch: ransomware detection through classification of high-entropy file segments,” *Journal of Cybersecurity*, vol. 11, no. 1, p. tyaf009, 2025.
- [8] C. Meijer, V. Moonsamy, and J. Wetzels, “Where’s crypto?: Automated identification and classification of proprietary cryptographic primitives in binary code,” in *Proc. 30th USENIX Security Symposium*, 2021.
- [9] G. Barnes and A. Ghafarian, “Machine learning-based static ransomware detection using PE header features and SHAP interpretation,” *Journal of Cybersecurity and Privacy*, vol. 6, no. 2, p. 58, 2026.
- [10] M. H. M. Baig, H. B. Ul Haq, and W. Habib, “A comparative analysis of AES, RSA, and 3DES encryption standards based on speed and performance,” *Management Science Advances*, vol. 1, no. 1, pp. 20–30, 2024.
- [11] Y. Banu, B. K. Rath, and D. Gountia, “Analyzing cryptographic algorithm efficiency within graph-based encryption models,” *Frontiers in Computer Science*, vol. 7, 2025, art. 1630222.
- [12] H. Oz, A. Aris, A. Levi, and A. S. Uluagac, “A survey on ransomware: Evolution, taxonomy, and defense solutions,” *ACM Computing Surveys*, vol. 54, no. 11s, 2022.
- [13] W. N. A. Ruzai, Y. Ying, K. N. Muhammad, M. A. Asbullah, and M. R. K. Ariffin, “Concurrent factorization of RSA moduli via weak key equations,” *AIMS Mathematics*, vol. 9, no. 10, pp. 28 211–28 231, 2024.
- [14] E. Pelofske, “An efficient all-to-all GCD algorithm for low entropy RSA key factorization,” *Cryptology ePrint Archive*, Paper 2024/699, 2024.
- [15] A. Raghu, “A context-aware framework for secure fintech APIs: Leveraging Java Spring Boot and OAuth 2.0 with dynamic token adaptation,” *Journal of Information Systems Engineering and Management*, vol. 10, no. 49s, pp. 367–377, 2025.
- [16] P. Rujichaikul and I. Rassameeroj, “Token-based authentication monitoring system,” *Journal of Cyber Security and Mobility*, vol. 14, no. 4, pp. 777–798, 2025.
- [17] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, “Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection,” in *Proc. 16th ACM Workshop on Artificial Intelligence and Security (AISec)*, 2023.
- [18] K. Hines, G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman, “Defending against indirect prompt injection attacks with spotlighting,” arXiv:2403.14720, 2024.
- [19] T. Gokcimen, “A novel system for strengthening security in large language models,” *Alexandria Engineering Journal*, vol. 75, pp. 100–112, Jan. 2025.
- [20] Y. Liu, Y. Jia, R. Geng, J. Jia, and N. Z. Gong, “Formalizing and benchmarking prompt injection attacks and defenses,” in *Proc. 33rd USENIX Security Symposium*, 2024.

APPENDIX A ADDITIONAL DESIGN DIAGRAMS

This appendix collects the remaining design diagrams: the level-1 data flow, the activity flow for grading injected code, the UML sequence of a defended run, and the main classes and modules.




Fig. 10. DFD level 1. Processes are circles, data stores are open rectangles, external entities are boxes.




Fig. 11. UML activity diagram: grading an injected-code region.




Fig. 12. UML sequence diagram of an analysis run with the defensed LLM report.




Fig. 13. UML class and module diagram (main classes and public functions).

APPENDIX B

LITERATURE SURVEY TABLE

TABLE VIII: Literature survey: qualitative relevance and reported cryptographic parameters of 20 works compared against TriageGuard

#	Author, Paper & Venue	Yr.	Domain Focus & Methodology	Key Findings & Observation	Key Size	PL Size	CT Size / Output	Encryption	Time Consuming	Performance	Limitation / Gap in Paper	Relevance: Memory Forensics	Relevance: Cryptographic Stress	Relevance: AI Robustness & Gap Addressed by TriageGuard
1	Baig, Ul Haq, Habib, "A Comparative Analysis of AES, RSA, and 3DES Encryption Standards based on Speed and Performance" (Management Science Advances) [10]	2024	Symmetric vs. asymmetric cipher overview; literature-based comparative review of AES, DES, 3DES, RSA covering computation time and memory usage	Symmetric ciphers process faster than RSA; RSA offers stronger key-exchange properties. A review rather than a fresh experimental benchmark; no primary data of its own	AES 128/192/256-bit; DES 56-bit; 3DES 112-bit	Not experimentally varied	N/A	AES, DES, 3DES, RSA	Cites prior benchmarks showing AES/DES faster than RSA	Qualitative time/memory comparison only	Relies entirely on cited studies; no primary experimental data	None directly	Confirms the baseline speed hierarchy (AES > DES > RSA) that TriageGuard's algorithm-classification module assumes	None directly
2	Banu, Rath, Gountia, "Analyzing cryptographic algorithm efficiency within graph-based encryption models" (Frontiers in Computer Science) [11]	2025	Asymmetric cipher performance on real files; implements RSA/ElGamal in C#, benchmarks on text/image/audio files of 22 KB to 5,120 KB against AES-256/ECC baselines	RSA is faster but memory-heavier; ElGamal is slower but more memory-efficient; both far slower than AES-256. Real measured trade-off; single-machine benchmark	RSA/ElGamal 2048-bit; AES-256; 64/128-bit in early trials	2,345 KB (text/image/ audio)	RSA ~169.82 KB (text/image/ audio); ElGamal ~0.17 KB	RSA, ElGamal vs. AES-256, ECC	RSA 0.108 s vs. ElGamal 1.55 s for 22 KB text; AES-256 ~0.001 s	RSA throughput ~203 KB/s vs. ElGamal ~14 KB/s	Single-machine benchmark; no malware-context evaluation	File-level, not memory-resident, but technique is reusable	Very high; supplies real operating ranges for TriageGuard's crypto-detection thresholds	None directly
3	Ruzai, Ying, Muhammad, Asbullah, Ariffin, "Concurrent factorization of RSA moduli via weak key equations" (AIMS Mathematics) [13]	2024	RSA cryptanalysis; applies Diophantine approximation and Coppersmith's lattice technique to sets of RSA public keys	Polynomial-time factorization achievable when small-integer relationships hold among RSA parameters. Purely theoretical; no implementation	RSA moduli, structurally related, unspecified bit-lengths	N/A	N/A	RSA	Polynomial-time (theoretical)	Demonstrates concurrent factorization of multiple weak moduli	No automated tooling to detect these conditions in real malware samples	None directly	Very high; underlies TriageGuard's RSA weak-key checks (small exponent, shared modulus)	None directly
4	Raghu, "A Context-Aware Framework for Secure Fintech APIs: Leveraging Java Spring Boot and OAuth 2.0 with Dynamic Token Adaptation" (J. Info. Systems Eng. & Mgmt.) [15]	2025	OAuth 2.0/JWT adaptive security for fintech APIs; Spring Boot + Redis-based dynamic risk scoring adjusting JWT scope/TTL in real time, load-tested with JMeter	Adaptive tokens cut unauthorized access by 84% at a small latency cost. Security gains carry an explicit, quantified performance cost	RSA 2048-bit (JWT via RS256)	N/A (tokens, not files)	N/A; JWT byte-size not reported	OAuth 2.0/JWT (RS256) with dynamic scope/TTL	+3 ms median latency (10 ms vs. 7 ms baseline)	≥230 TPS vs. 250 baseline; 84% fewer unauthorized attempts	Tested on one GKE deployment; not evaluated against forged tokens embedded in malware	None directly	High; benchmark for acceptable JWT/OAuth hardening overhead	None directly
5	Rujichaikul, Rassameeroj, "Token-Based Authentication Monitoring System" (J. Cyber Security and Mobility) [16]	2025	JWT/OAuth attack detection; 25 rule-based detection signatures tested via 70 test cases on a token-based web app	81.4% detection accuracy, 92% recall, low operational overhead. Prioritises recall over false-positive rate	N/A; monitors usage patterns, not the signing key	N/A	N/A	JWT access tokens + OAuth 2.0 refresh tokens (monitored)	9.7 ms avg. detection latency	81.4% accuracy; 2.3% CPU overhead; +18 MB memory	Struggles with VPN/IP-spoofing edge cases; no ML component	None directly	High; precedent for detection-accuracy/overhead trade-offs	None directly
6	Oh, Kim, Kim, Kim, "volGPT: Evaluation on triaging ransomware process in memory forensics with large language model" (Forensic Sci. Int.: Digital Investigation) [2]	2024	AI-assisted memory forensics for ransomware; extracts process artefacts from memory images, feeds them to an LLM to triage ransomware-related processes	LLM-assisted triage improves identification of ransomware-relevant processes over raw plugin output. Prompt design and hallucination remain open issues	N/A	N/A	N/A	N/A (detects ransomware presence, not cipher parameters)	Not reported	Improved triage accuracy vs. plugin baseline	Never tests whether a crafted sample can bias the LLM's own verdict	Very high; near-identical Stage 1 approach TriageGuard builds on	Low; detects ransomware but not encryption strength	Gap TriageGuard's red-teaming track directly addresses

#	Author, Paper & Venue	Yr.	Domain Focus & Methodology	Key Findings & Observation	Key Size	PL Size	CT Size / Output	Encryption	Time Consuming	Performance	Limitation / Gap in Paper	Relevance: Memory Forensics	Relevance: Cryptographic Stress	Relevance: AI Robustness & Gap Addressed by TriageGuard
7	Sharma, Ghawaly, McCleary, Webb, Baggili, "ForensicLLM: A local large language model for digital forensics" (Forensic Sci. Int.: Digital Investigation) [3]	2025	Locally-hosted LLM for forensic reports; fine-tunes a local LLM on forensic artefact data for investigator queries without cloud exposure	Local deployment matches cloud LLM Q&A performance with less data-exposure risk. Focuses on deployment/privacy, not adversarial evidence	N/A	N/A	N/A	N/A	Not reported	Comparable to cloud LLM performance	No test of manipulated input	High; validates TriageGuard's own LLM-report design choice	None directly	Adds the missing layer TriageGuard tests: resistance to manipulated input
8	Casino, Hurley-Smith, Hernandez-Castro, Patsakis, "Not on my watch: ransomware detection through classification of high-entropy file segments" (Journal of Cybersecurity) [7]	2025	Entropy-based ransomware/encryption detection; trains a classifier to distinguish genuinely encrypted segments from other high-entropy data (e.g. compressed files)	Outperforms naive entropy-threshold detectors, reduces false positives. Detects that data is encrypted, not which algorithm or how strong	N/A (detection, not a specific cipher benchmark)	File segments of varying entropy (size unspecified)	N/A	Detects encrypted vs. compressed high-entropy data (algorithm-agnostic)	Not reported	Improved false-positive rate over threshold baseline	No algorithm identification or key-strength evaluation	Moderate; file-level, technique reusable	Very high; direct basis for TriageGuard's entropy module	None; no AI reporting layer
9	Gokcimen, "A novel system for strengthening security in large language models" (Alexandria Engineering Journal) [19]	2025	LLM security hardening against manipulation; proposes an architectural defence layer for LLM deployments, evaluated against injection-style attacks	Measurable improvement in resistance to malicious input manipulation. General-purpose framework, not forensics-specific	N/A	N/A	N/A	N/A	Not reported	Reduced attack success rate (framework-level)	Not evaluated in a forensics/malware-triage context or against evidence-embedded attacks	None directly	None directly	Very high; parallels TriageGuard's own attack-then-harden methodology
10	Barnes, Ghafarian, "Machine Learning-Based Static Ransomware Detection Using PE Header Features and SHAP Interpretation" (J. Cybersecurity and Privacy) [9]	2026	Static PE-header-based ransomware classification; trains Random Forest, SVM, XGBoost on PE header features across three datasets; uses SHAP to interpret decisions	Strong classification accuracy; SHAP identifies the most indicative header fields. Static features are brittle against packing/obfuscation	N/A	N/A	N/A	N/A (static feature classification, not cipher benchmark)	Not reported	High classification accuracy (exact figures not in abstract)	Doesn't test evasion or combine with entropy/behavioural signals	High; reinforces TriageGuard's static-analysis	Moderate; complements entropy/key-level analysis	None; no adversarial-robustness component
11	Carrier, Victor, Tekeoglu, Lashkari, "Detecting Obfuscated Malware using Memory Feature Engineering" (ICISSP) [1]	2022	ML malware detection from memory; extends the VolMemLyzer feature extractor (built on Volatility) and trains a stacked ensemble on the new MalMemAnalysis-2022 dataset	Memory features detect obfuscated malware with 99.00% accuracy and 99.02% F1-score	N/A	N/A	N/A	N/A	Not reported here	99.00% accuracy, 99.02% F1	Binary malicious/benign verdict; no explanation, no cryptographic analysis	Very high; confirms Volatility features carry the evidence TriageGuard's Stage 1 extracts	None directly	None; no LLM layer. TriageGuard explains findings in a cited report instead of a label
12	Scanlon, Breitingner, Hargreaves, Hilgert, Sheppard, "ChatGPT for digital forensic investigation: The good, the bad, and the unknown" (Forensic Sci. Int.: Digital Investigation) [4]	2023	LLMs in forensics; tests GPT-4 on artefact analysis, evidence search, script writing, anomaly detection, incident response and training scenarios	Some low-risk uses work, but many are unsuitable today; uploading evidence raises privacy issues and the user must catch the model's mistakes	N/A	N/A	N/A	N/A	Not reported	Qualitative, per scenario	No systematic defence against wrong or manipulated output	High; motivates feeding the LLM findings, not raw evidence	None directly	Very high; TriageGuard adds mandatory citations and a rule cross-check so mistakes are caught automatically
13	Fernandez de Loaysa Babiano, Macfarlane, Davies, "Evaluation of live forensic techniques, towards Salsa20-based cryptographic ransomware mitigation" (Forensic Sci. Int.: Digital Investigation) [5]	2023	Live memory forensics against ransomware; extracts Salsa20 keys and nonces from volatile memory while ransomware runs	Over 90% of Salsa20 key and nonce pairs recovered from memory, enough to decrypt files without paying, even with a unique key per file	Salsa20 256-bit key + 64-bit nonce	Victim files	Encrypted files (stream cipher, same length)	Salsa20 inside hybrid encryption	Not reported here	>90% key/nonce pairs recovered	Targets Salsa20 specifically; needs capture while the ransomware is running	Very high; shows modern ransomware keys live in RAM	Very high; supports TriageGuard's ChaCha/Salsa20 detection	None

#	Author, Paper & Venue	Yr.	Domain Focus & Methodology	Key Findings & Observation	Key Size	PL Size	CT Size / Output	Encryption	Time Consuming	Performance	Limitation / Gap in Paper	Relevance: Memory Forensics	Relevance: Cryptographic Stress	Relevance: AI Robustness & Gap Addressed by TriageGuard
14	Baier, Basse, Hilgert, Lambertz, "TLS key material identification and extraction in memory: Current state and future challenges" (DFRWS USA) [6]	2024	Systematisation of knowledge; reviews academic and in-the-wild methods for finding and extracting TLS key material from memory	Memory extraction of TLS keys is essential for decrypting malware traffic; lists open research challenges	TLS session keys and secrets	N/A	N/A	TLS (symmetric session keys)	Not reported	Qualitative survey	Survey; does not cover RSA weak-key testing or reporting	Very high; memory as a source of key material	High; links TriageGuard's memory and pcap stages	None
15	Meijer, Moonsamy, Wetzels, "Where's Crypto?: Automated Identification and Classification of Proprietary Cryptographic Primitives in Binary Code" (USENIX Security) [8]	2021	Crypto identification in binaries; data-flow-graph isomorphism plus symbolic execution, released as an IDA plugin	Identifies and classifies both known and proprietary primitives, tested on public and proprietary binaries	Varies per primitive	N/A	N/A	Known and proprietary primitives	Not reported here	Not reported here	Per-binary analysis in a disassembler; not on whole memory images	Low	Very high; covers custom ciphers that TriageGuard's constant signatures would miss	None
16	Oz, Aris, Levi, Uluagac, "A Survey on Ransomware: Evolution, Taxonomy, and Defense Solutions" (ACM Computing Surveys) [12]	2022	Survey of ransomware evolution, taxonomy, and detection and defence techniques across platforms	Organises ransomware by target, infection and encryption behaviour, and reviews defences	Varies (symmetric + asymmetric hybrids)	Victim files	Encrypted files	Hybrid encryption, varies by family	Not reported	Survey	Survey; no tool	Moderate	High; motivates identifying which cipher a sample really uses	None
17	Pelofske, "An Efficient All-to-All GCD Algorithm for Low Entropy RSA Key Factorization" (IACR ePrint 2024/699) [14]	2024	RSA weak-key cryptanalysis; binary-tree batch GCD to find RSA moduli that share a prime because of low-entropy key generation	About 6x faster than the standard remainder-tree batch GCD, same asymptotic scaling	RSA moduli	N/A	Shared prime factors	RSA	About 6x speedup over remainder-tree batch GCD	Speedup grows with the number of shared factors	Algorithmic; applied to key collections, not memory	None directly	Very high; same shared-prime test TriageGuard runs on carved keys	None
18	Greshake, Abdelnabi, Mishra, Endres, Holz, Fritz, "Not What You've Signed Up For: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection" (ACM AISec) [17]	2023	Indirect prompt injection; instructions planted in data an LLM application retrieves, shown on real systems	Retrieved data can take over an LLM-integrated application	N/A	N/A	N/A	N/A	Not reported	Demonstrated on deployed applications	General applications; no forensic setting; no measured defences	Low	None directly	Very high; memory artefacts are exactly such retrieved data
19	Hines, Lopez, Hall, Zarfati, Zunger, Kiciman, "Defending Against Indirect Prompt Injection Attacks With Spotlighting" (arXiv:2403.14720) [18]	2024	Defence; marks untrusted input by delimiting, datamarking or encoding so the model can tell data from instructions	Reduces attack success from over 50% to under 2% in their experiments	N/A	N/A	N/A	N/A	Not reported	ASR from >50% to <2%	Single input-side defence; no output validation	None directly	None directly	Very high; TriageGuard's defence D2, combined with output checks
20	Liu, Jia, Geng, Jia, Gong, "Formalizing and Benchmarking Prompt Injection Attacks and Defenses" (USENIX Security) [20]	2024	Framework formalising prompt injection; benchmark of 5 attacks and 10 defences on 10 LLMs and 7 tasks	Combined attacks are most effective; existing defences are not sufficient on their own	N/A	N/A	N/A	N/A	Not reported	Benchmark across 10 LLMs	General tasks; no evidence-citation check	None directly	None directly	Very high; supports layering several defences as TriageGuard does